

A Crash Course in Python Performance

From identifying slow code, to calling c libraries from your python.

by Tommi Kabelitz (University of Adelaide)

on 25/03/2022

- * `git@github.com:TommiKabelitz/FastPython.git`
- * `https://github.com/TommiKabelitz/FastPython`

Profiling

» Identifying slow code

INVEST YOUR TIME WHERE YOUR CODE TAKES TIME

- * You may think you know what is slow - You will often be wrong
- * Profilers track the calls to functions, keeping track of how long is spent where
- * Reports are (normally) easy to read and sort

» Identifying slow code

- * Note: slight overhead associated with using profilers
- * Use at the start to work out which functions are slow
- * When improving functions, use modules like timeit to simply time the function without overhead
- * I recommend cProfile for python

» Using the profiler

Module imports

```
#The profiler
import pstats, cProfile

#For tidying the output
import io
from pstats import SortKey
```

» Using the profiler

Running the profiler

#Initialising the profiler

```
prof = cProfile.Profile()
```

#Profiling some function

```
prof.enable()
```

```
function_to_profile()
```

```
prof.disable()
```

» Using the profiler

Outputting the stats

```
#A little messing around is required to format the output
s = io.StringIO()
sortby = SortKey.CUMULATIVE #Alternatively TIME, CALLS, etc
ps = pstats.Stats(prof, stream=s).sort_stats(sortby)
ps.print_stats(100)
#Now the output is in s, so either
print(s.getvalue())
#or
with open('profile.txt', 'w') as f:
    f.write(s.getvalue())
```

» Using the profiler

Profiler output (Example from my plotting code)

42441171 function calls (40411093 primitive calls) in 30.739 seconds

Ordered by: cumulative time

List reduced from 2678 to 100 due to restriction <100>

	ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
	1	0.000	0.000	30.767	30.767	effectiveMass.py:21(main)
	1	0.018	0.018	30.750	30.750	effectiveMass.py:68(DoCombination)
	254	0.001	0.000	24.578	0.097	/*condapath*/matplotlib/backends/backend_pdf.py:2464(savefig)
	254	0.001	0.000	24.545	0.097	/*condapath*/matplotlib/figure.py:2063(savefig)
	254	0.005	0.000	24.543	0.097	/*condapath*/matplotlib/backend_bases.py:2001(print_figure)
	254	0.003	0.000	24.444	0.096	/*condapath*/matplotlib/backends/backend_pdf.py:2532(print_pdf)
35808/254	0.088	0.000	0.000	24.352	0.096	/*condapath*/matplotlib/artist.py:30(draw_wrapper)
	254	0.006	0.000	24.351	0.096	/*condapath*/matplotlib/figure.py:1688(draw)
	256	0.001	0.000	12.555	0.049	/*condapath*/matplotlib/cbook/deprecation.py:347(wrapper)
	256	0.004	0.000	12.547	0.049	/*condapath*/matplotlib/figure.py:2448(tight_layout)
	256	0.004	0.000	12.430	0.049	/*condapath*/matplotlib/tight_layout.py:264(get_tight_layout_fig
ure)						
	256	0.010	0.000	12.384	0.048	/*condapath*/matplotlib/tight_layout.py:33(auto_adjust_subplotpa
rs)						
	256	0.001	0.000	12.310	0.048	/*condapath*/matplotlib/tight_layout.py:109(<listcomp>)
	256	0.037	0.000	12.309	0.048	/*condapath*/matplotlib/axes/_base.py:4270(get_tightbbox)
	508/254	0.007	0.000	11.743	0.046	/*condapath*/matplotlib/image.py:119(_draw_list_compositing_imag
es)						

» Using the profiler

Columns

- * ncalls: Number of function calls. Denominator of fraction (if present) denotes number of recursive calls
- * tottime: Time spent in a function, excluding time spent in other functions
- * percall: $\text{tottime} / \text{ncalls}$
- * cumtime: Total time spent in a function, including calls to subfunctions. Accurate for recursive functions
- * percall: $\text{cumtime} / \text{ncalls}$
- * filename:lineno(function) The function to which the stats refer

Profiling
oooooooo●●

Weaknesses
ooo

Leverage other languages
oooo

Vectorisation
oooo

MPI
oo

Cython
oooooooooooooooo

» Using the profiler

Takeaways about my plotting code

- * Majority of time is spent saving figures to file
- * Time reading/manipulating data is minimal
- * If I want to speed up this code (I don't, not worth it)
 - * Save speed is only important factor. Ideas:
 - * Plots per page
 - * Maybe pdf is not ideal
 - * Different package?
 - * etc.

Profiling
ooooooooo●

Weaknesses
ooo

Leverage other languages
oooo

Vectorisation
oooo

MPI
oo

Cython
oooooooooooooooooooo

» Profiling recap

- * When you care about speed, profile. It is easy
- * You will be surprised at which piece of code is the slow part
- * You invest your time improving the code that is slow

Weaknesses

» A word on why python can be slow

We cannot write efficient python without knowing what makes it slow.
Hence, weaknesses:

- * Interpreted not compiled
- * Duck typing
- * The GIL (Global Interpreter Lock)

» A word on why python can be slow

A simple example

```
n = 10
scale = 0.4
if scale > 1:
    total = 0
    for i in range(1,n+1):
        total += scale*i
else:
    total = 1
    for i in range(1,n+1):
        total *= (scale - i)
```

Leverage other languages

Profiling
oooooooooooo

Weaknesses
ooo

Leverage other languages
o●oo

Vectorisation
oooo

MPI
oo

Cython
oooooooooooooooooooo

» Integrating other languages

- * Why python? → easy to write
- * But we tradeoff speed for that
- * Consider mixing languages

» Integrating other languages

- * Do the heavy computation in c, c++, fortran, etc.
- * Do the fiddly setup in python
 - * File paths
 - * String manipulation
 - * Data cleaning
 - * Parameter organisation
 - * etc.

» Integrating other languages

The subprocess module is excellent for running other code. A simple example:

```
#the command to run, each argument as an element
executable = ['path/to/executable.x']
report_file = 'path/to/reportfile.txt'
input_file = 'path/to/input/file.txt'

generate_input_file(input_file) #Generate your input file(s)

with open(report_file, 'w') as f:
    subprocess.run(command,
                    input=input_file + '\n',
                    text=True,
                    stdout=f,
                    stderr=subprocess.STDOUT,
                    timeout=600) #timeout in seconds

# Note: don't just copy paste this entire code block,
# the indentation is broken because Tex hates me.
```

Vectorisation

» Writing *fast* pure-python

- * Cannot expect speed in any language if the code is *bad*
- * Vectorisation is vital
- * Many python libraries provide highly vectorised routines already

Vectorisation

The transformation of code to act on entire arrays at once, rather than element by element

» Writing *fast* pure-python

Vectorisation steps

- * Remove loops where possible
- * Replace loops with vectorised operations

```
import numpy as np
n = 10
scale = 0.4
array = np.arange(1,n+1)
if scale > 1:
    total = scale*sum(array)
else
    total = np.prod(scale - array)
```

Profiling
oooooooooooo

Weaknesses
ooo

Leverage other languages
oooo

Vectorisation
ooo●

MPI
oo

Cython
oooooooooooooooooooo

» Writing *fast* pure-python

- * numpy, scipy and libraries written on those are likely to be highly vectorised
- * So use functions from there as often as possible
- * There are a lot
- * But sometimes creativity is required

MPI

» Leveraging multiple cores

- * Very possible in python
- * Several libraries for it
- * Definitely something to consider when you do not need any message passing
- * Still possible with message passing, but more painful

DISCLAIMER: I haven't ever used multiple CPUs in python as I haven't needed to. I just want to mention it.

Cython

- * A basic intro
 - * For directly integrating other languages
 - * For turning python code into c code

Profiling
oooooooooooo

Weaknesses
ooo

Leverage other languages
oooo

Vectorisation
oooo

MPI
oo

Cython
o●oooooooooooo

» What is cython

- * Compiles python into c (or c++)
- * Functions can be directly imported into python
- * Can import c and fortran functions too
- * Basics are very simple, plenty of complex stuff too

Profiling
oooooooooooo

Weaknesses
ooo

Leverage other languages
oooo

Vectorisation
oooo

MPI
oo

Cython
oo●oooooooooooo

» Installation

Installation:

```
python -m pip install cython
```

- * File extension is .pyx
- * Compilation will produce a .so file
- * Then simply import as normal

» Compilation

Compile using

```
python setup.py build_ext --inplace
```

Where the setup.py file is

```
from setuptools import setup
from Cython.Build import cythonize

setup(ext_modules=cythonize('cython_module.pyx'),
      compiler_directives={'language_level':3},
      zip_safe=False)
```

» The .pyx file

In terms of preparing the .pyx file all you need is:

```
import cython
```

```
# Optional compiler directives
```

```
# cython: boundscheck=False, wraparound=False
```

```
# cython: initializedcheck=False
```

Cython

- * A basic intro
 - * For directly integrating other languages
 - * For turning python code into c code

» Definitions

The most important pieces of cython syntax:

- * cdef - A purely c, fastest
- * cpdef - A hybrid object, slower, importable in python

Rule of thumb, if you don't need to use/access the value from within your python. Use cdef.

Only cpdef the functions/objects you need at the end

» Typing

You do not have to type everything. The more you type, the faster your code.

#Function definition (returns a 128bit int)

```
cpdef long my_fun(int a, float b,  
                  str s, double[:] array):
```

```
    cdef int i,j,k  
    cdef double* array_pointer = &array[0]  
    cdef int size = array.shape[0]  
    c = 2*a      #Note not typed, but could be
```


Cython

- * A basic intro
 - * For directly integrating other languages
 - * For turning python code into c code

Profiling
oooooooooooo

Weaknesses
ooo

Leverage other languages
oooo

Vectorisation
oooo

MPI
oo

Cython
oooooooooooo●oooo

» Cython for integrating other languages

- * subprocess.run is not bad
- * Still awkward passing information between languages
- * Cython allows direct calling of c, c++ and fortran subroutines
- * And it actually is pretty simple

» Cython for integrating other languages

A little involved for slides, so see the github repo in `cython/interfacing/fortran`. Files required: (and the order to look at them)

- * `fortmod.f90` (Fortran module containing code to import)
- * `fort_interface.f90` (Interfacing the fortran to c)
- * `fort_interface.h` (c header file)
- * `fort_interface.pyx` (cython file which interfaces the python and fortran)
- * `setup.py` (Has to change to link the code properly)
- * `runscript.py` (Does the actual call)

Cython

- * A basic intro
 - * For directly integrating other languages
 - * For turning python code into c code

» What python should I compile with cython?

Compiling all of your python is a waste of time. Have to recompile every time, cython only helps with computation speed.

Compile:

- * The functions that take lots of time in your profile report
- * (Only if typing is the bottleneck)
- * Functions that get called often
- * Functions which cannot be vectorised easily
- * Functions which you don't touch often

Examples:

- * Jackknife/Bootstrap routines
- * Functions that are passed to optimisation routines

» Abrupt ending

Takeaways

- * Please profile your code
- * Various options for speedup (vectorisation, mpi, interfacing other languages)
- * Cython is also an option
- * Either use it to type your code
- * IMO its true calling is direct interfacing

Questions?